

# Software Project: Image Compression using SVD

ee25btech11009 – Anshu kumar ram

## Contents

<b>1 Overview and SVD basics</b>	<b>2</b>
<b>2 SVD for image compression</b>	<b>2</b>
2.1 Mathematical formulation . . . . .	2
<b>3 Implementation : power iteration + deflation</b>	<b>3</b>
3.1 Pseudocode (implementation-accurate) . . . . .	3
<b>4 Mathematical concept for Power Iteration</b>	<b>4</b>
<b>5 Error metric</b>	<b>4</b>
<b>6 Why power iteration was chosen</b>	<b>4</b>
6.1 Why some alternatives were not used . . . . .	5
<b>7 Practical choices: iterations and stability</b>	<b>5</b>
<b>8 Observations, limitations and trade-offs</b>	<b>5</b>
8.1 Rank $k$ vs compression . . . . .	5
<b>9 Results (example tables and images)</b>	<b>7</b>
<b>10 Conclusion</b>	<b>12</b>

# 1 Overview and SVD basics

Summary of Strang's lecture.

Any real matrix  $A$  admits a factorization

$$A = U\Sigma V^\top, \quad (1)$$

where  $U^\top U = I$ ,  $V^\top V = I$ , and  $\Sigma$  is diagonal with nonnegative entries:

$$\Sigma = \begin{pmatrix} \sigma_1 & 0 & \cdots & 0 \\ 0 & \sigma_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_r \end{pmatrix}, \quad \sigma_1 \geq \sigma_2 \geq \cdots \geq \sigma_r > 0.$$

The  $\sigma_i$  are the singular values of  $A$ .

Conceptually:

- The columns of  $V$  provide an orthonormal basis for the row space (right singular vectors).
- The columns of  $U$  provide an orthonormal basis for the column space (left singular vectors).
- Each right singular vector  $v_i$  is mapped by  $A$  to  $\sigma_i u_i$ :

$$Av_i = \sigma_i u_i.$$

To compute these components one inspects

$$A^\top A = V\Sigma^\top \Sigma V^\top, \quad (2)$$

$$AA^\top = U\Sigma\Sigma^\top U^\top, \quad (3)$$

so the eigenvectors of  $A^\top A$  are the columns of  $V$  with eigenvalues  $\sigma_i^2$ , and similarly for  $U$ .

A special case occurs when  $A$  is symmetric positive definite: then  $U = V$  and the SVD reduces to the eigendecomposition  $A = Q\Lambda Q^\top$  with positive eigenvalues.

## 2 SVD for image compression

A grayscale image of size  $m \times n$  can be encoded as a matrix  $A \in \mathbb{R}^{m \times n}$ , where each entry represents the pixel intensity (typically in  $[0, 255]$ ). The aim of truncated SVD compression is to approximate  $A$  by a lower-rank matrix  $A_k$  that captures most of the visually-important structure:

$$A \approx A_k = \sum_{i=1}^k \sigma_i u_i v_i^\top. \quad (4)$$

Using only the top  $k$  singular triplets reduces storage while preserving the principal image features.

### 2.1 Mathematical formulation

Let  $r = \text{rank}(A)$ . The full SVD is  $A = U\Sigma V^\top$ . The optimal (in Frobenius norm) rank- $k$  approximation is

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^\top. \quad (5)$$

This minimizes  $\|A - A_k\|_F$  among all matrices of rank at most  $k$ .

### 3 Implementation : power iteration + deflation

**Function:** `compute_svd()`

This implementation computes the top- $k$  singular triplets of a matrix  $A$  using power iteration applied to the symmetric matrix  $A^\top A$ . (compute  $A^\top A$ , extract eigenpairs with power iteration and deflation on  $A^\top A$ , then form  $U$  from  $A$ ,  $V$ , and the singular values).

The main steps are:

- Compute the transpose  $A^\top$  and form the symmetric Gram matrix  $A^\top A$ .
- Allocate storage for the top  $k$  eigenvalues and the matrix of right singular vectors  $V \in \mathbb{R}^{n \times k}$  (where  $n$  is the number of columns of  $A$ ).
- For  $i = 1 \dots k$  do:
  - Use the Power Iteration method on  $A^\top A$  to approximate the  $i^{\text{th}}$  dominant eigenvalue  $\lambda_i$  and eigenvector  $v_i$ .
  - Store  $v_i$  as the  $i^{\text{th}}$  column of  $V$ .
  - Deflate the matrix  $A^\top A$  by removing the contribution of the found eigenpair so the next iteration recovers the next eigenvector.
- Convert eigenvalues to singular values via  $\sigma_i = \sqrt{|\lambda_i|}$ . Taking absolute value is a defensive choice against tiny negative round-off errors.
- Compute the left singular vectors  $U$  from  $A$ ,  $V$  and  $\Sigma$  using

$$U = AV\Sigma^{-1},$$

i.e. compute each column  $u_i = \frac{Av_i}{\sigma_i}$ . Columns corresponding to  $\sigma_i \approx 0$  must be handled carefully (skip or orthonormalize).

- Pack the results into an SVD structure (or equivalent) containing  $U$ ,  $\sigma$ ,  $V$ , and  $k$ .
- Free temporary matrices and arrays (e.g.,  $A^\top$ ,  $A^\top A$ , intermediate eigenvalue arrays) to avoid memory leaks.

These steps match the provided implementation where power iteration and deflation are performed on  $A^\top A$ , and  $U$  is computed afterward from  $A$  and the computed  $V, \Sigma$ .

#### 3.1 Pseudocode (implementation-accurate)

```

Input: image file -> matrix A (m x n)
Compute AT = transpose(A)
Compute ATA = AT * A    # symmetric n x n matrix
Allocate arrays: eigenvalues[1..k], V[n x k]
for i = 1..k:
  v = random nonzero vector in R^n
  normalize v
  repeat for max_iteration:
    v = ATA * v          # power iteration on ATA
    normalize v
  end repeat
  lambda = v^T * (ATA * v)  # Rayleigh quotient (approx eigenvalue)
  store lambda in eigenvalues[i]
```

```

store v as column i of V
if i < k:
    ATA = ATA - lambda * (v * v^T)    # deflation on ATA
end for
sigma[i] = sqrt(abs(eigenvalues[i])) for i=1..k
Compute U: for i=1..k: u_i = (A * v_i) / sigma[i]
Package SVD = {U, sigma, V, k}
Free temporary memory

```

Write compressed representation using top-k triplets

## 4 Mathematical concept for Power Iteration

Let  $M$  be a matrix (here  $M = A^\top A$ ). Power iteration finds the eigenvector corresponding to the largest-in-magnitude eigenvalue of  $M$ . Starting from an initial vector  $v_0$  with a nonzero component along the dominant eigenvector, repeated multiplication by  $M$  amplifies the dominant mode.

Assume an eigendecomposition  $M = X\Lambda X^{-1}$  with  $|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_n|$ . Writing

$$v_0 = \sum_j c_j x_j,$$

we have

$$v_k = M^k v_0 = \sum_j c_j \lambda_j^k x_j.$$

Dividing by  $\|v_k\|$  isolates the  $x_1$  direction as  $k$  grows. Normalization each step prevents overflow and stabilizes convergence.

## 5 Error metric

We measure error by the Frobenius norm:

$$E = \|A - A_k\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n (A_{ij} - (A_k)_{ij})^2}. \quad (6)$$

## 6 Why power iteration was chosen

The implementation favors power iteration because:

- It is straightforward to implement: only repeated matrix-vector multiplications and scalings are needed.
- It has a low memory footprint since it stores only a few vectors rather than full orthogonal bases.
- It is efficient when one needs only the top  $k$  singular components instead of the entire SVD.

## 6.1 Why some alternatives were not used

**Jacobi method:** converges slowly for large matrices, needs multiple sweeps, and scales poorly (roughly  $O(n^3)$ ).

**Golub–Kahan bidiagonalization:** offers robust convergence but requires more complex orthogonal transformations and storage of basis vectors.

**Divide-and-conquer SVD:** involves recursive decomposition and is both programmatically and memory intensive for general images.

## 7 Practical choices: iterations and stability

In the implementation each singular vector is refined using a fixed number of power iterations (e.g., 50). This fixed cap balances runtime and accuracy: for many images convergence occurs in fewer iterations, but the upper limit ensures stability even for ill-conditioned problems.

Normalization at each iteration is essential to prevent overflow/underflow and to stabilize convergence:

$$v \leftarrow \frac{v}{\|v\|}.$$

## 8 Observations, limitations and trade-offs

- **Convergence:** power iteration converges quickly when the dominant singular value is well separated; convergence slows if leading singular values are close.
- **Numerical effects:** normalization mitigates growth/decay of iterates, but round-off may still accumulate for very large matrices.
- **Perceptual vs numeric trade-off:** beyond a certain  $k$ , additional singular components yield diminishing perceptual benefits while increasing storage and time.

### 8.1 Rank $k$ vs compression

The storage needed to represent  $A_k$  using the compact SVD form is

$$N_k = k(m + n + 1),$$

while the original requires  $N_{\text{orig}} = mn$  scalars. The compression ratio is therefore

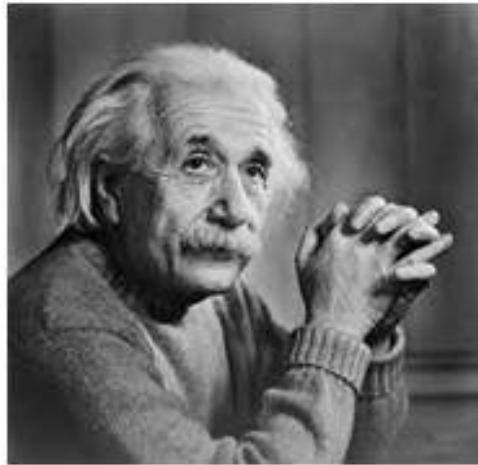
$$\text{Compression ratio} = \frac{mn}{k(m + n + 1)}.$$



## 9 Results (example tables and images)

Images

(a) Original Einstein Image

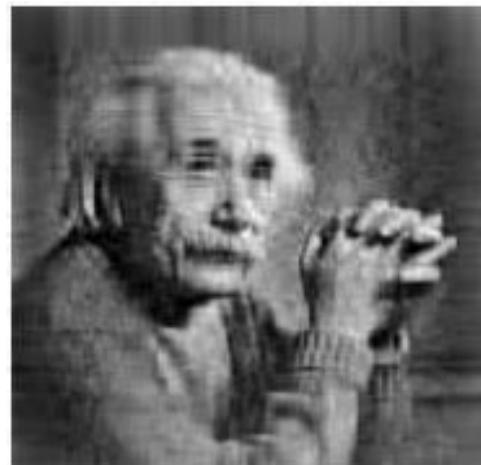


(a) Original Image

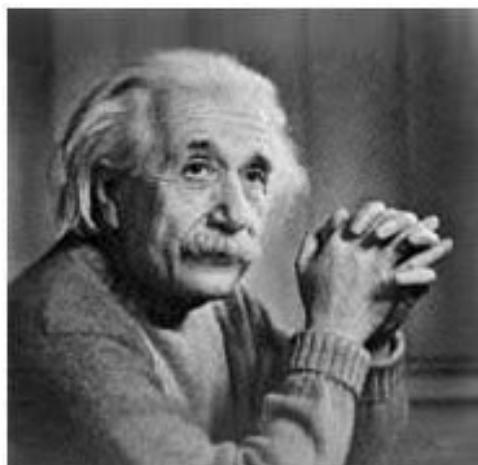
Compressed Images



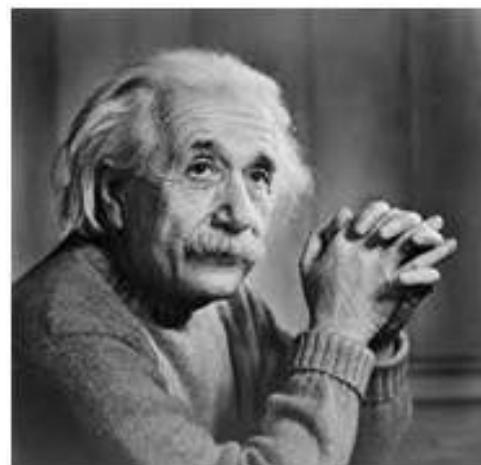
(b)  $k = 5$



(c)  $k = 20$



(d)  $k = 50$



(e)  $k = 100$

<b><math>k</math></b>	<b>Compression Ratio</b>	<b>Runtime (s)</b>	<b>Frobenius Error</b>
5	18.35	0.146	4713.5869
20	4.59	0.520	2126.5615
50	1.83	1.366	880.5452
100	0.92	2.735	165.0320

Table 1: SVD-based image compression performance for different values of  $k$  of einstein.jpg.

**Images**  
**(a) Original Globe Image**



**(a) Original Image**

**Compressed Images**



**(b)  $k = 5$**



**(c)  $k = 20$**



**(d)  $k = 50$**

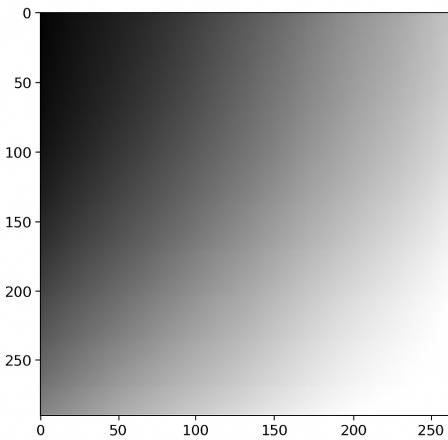


**(e)  $k = 100$**

<b><math>k</math></b>	<b>Compression Ratio</b>	<b>Runtime (s)</b>	<b>Frobenius Error</b>
5	85.86	9.158	20704.2746
20	21.46	18.604	10634.4134
50	8.59	36.708	6186.0730
100	4.29	63.059	3673.9582

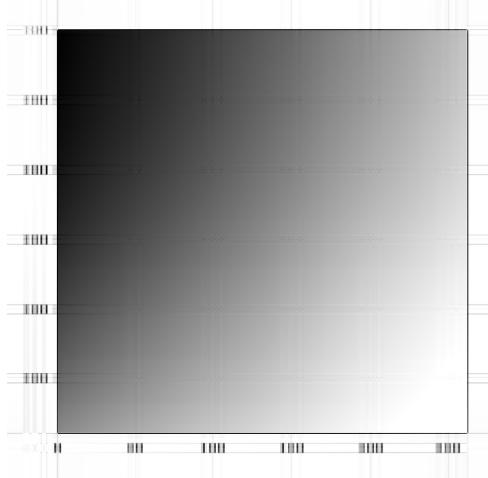
Table 2: SVD-based image compression performance for the `globe.jpg` image.

**Images**  
**(a) Original Greyscale Image**

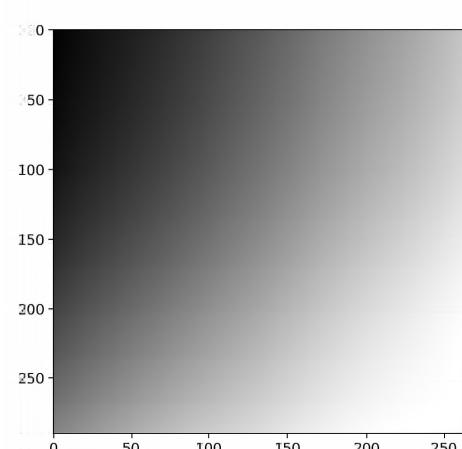


(a) Original Image

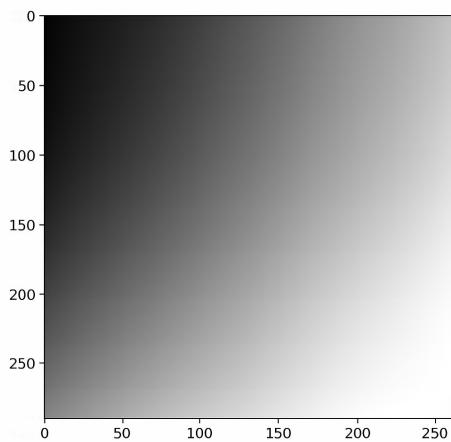
**Compressed Images**



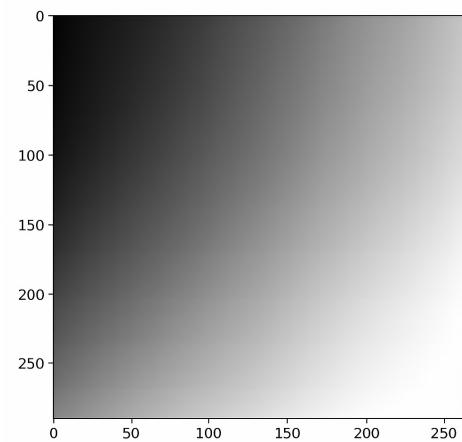
(b)  $k = 5$



(c)  $k = 20$



(d)  $k = 50$



(e)  $k = 100$

Figure 3: SVD-based image compression results for the greyscale image using different values of  $k$ .

<b><math>k</math></b>	<b>Compression Ratio</b>	<b>Runtime (s)</b>	<b>Frobenius Error</b>
5	102.35	17.570	11146.3094
20	25.59	28.696	3808.1952
50	10.24	56.157	1160.1419
100	5.12	100.557	512.4441

Table 3: SVD-based image compression performance for the `greyscale.png` image.

## 10 Conclusion

Truncated SVD offers an effective way to compress grayscale images: storing the top  $k$  singular triplets typically captures the bulk of perceptual detail with far fewer numbers than the full image. Power iteration with deflation is an accessible, memory-light strategy for extracting these dominant components when a full SVD is not necessary.